

IMPORTING LIBRARIES

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import datetime
from pylab import rcParams
import matplotlib.pyplot as plt
import warnings
import itertools
import statsmodels.api as sm
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import LSTM
from keras.layers import Dropout
from sklearn.metrics import mean_squared_error
from keras.callbacks import ReduceLROnPlateau, EarlyStopping, ModelCheckpoint
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
import seaborn as sns
sns.set_context("paper", font_scale=1.3)
sns.set_style('white')
import math
from sklearn.preprocessing import MinMaxScaler
# Input data files are available in the "../input/" directory.
# For example, running this (by clicking run or pressing Shift+Enter) will
# list all files under the input directory
warnings.filterwarnings("ignore")
plt.style.use('fivethirtyeight')
import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

IMPORTING DATA

```
import io
df = pd.read_excel('Crude Oil Prices Daily.xlsx')
df.head()
df[:10]
```

In [10]:

Out[10]:

| | Date | Closing Value |
|---|------------|---------------|
| 0 | 1986-01-02 | 25.56 |
| 1 | 1986-01-03 | 26.00 |
| 2 | 1986-01-06 | 26.53 |

| | Date | Closing Value |
|----------|-------------|----------------------|
| 3 | 1986-01-07 | 25.85 |
| 4 | 1986-01-08 | 25.87 |
| 5 | 1986-01-09 | 26.03 |
| 6 | 1986-01-10 | 25.65 |
| 7 | 1986-01-13 | 25.08 |
| 8 | 1986-01-14 | 24.97 |
| 9 | 1986-01-15 | 25.18 |

In [11]:

```
#Sort dataset by column Date
df = df.sort_values('Date')
df = df.groupby('Date')['Closing Value'].sum().reset_index()
df.set_index('Date', inplace=True)
df=df.loc[datetime.date(year=2000,month=1,day=1):]
```

In [12]:

```
df.head()
```

Out[12]:

| | Closing Value |
|-------------------|----------------------|
| Date | |
| 2000-01-04 | 25.56 |
| 2000-01-05 | 24.65 |
| 2000-01-06 | 24.79 |
| 2000-01-07 | 24.79 |

Closing Value

Date

2000-01-10 24.71

DATA PRE-PROCESSING

In [13]:

```
def DfInfo(df_initial):  
    # gives some infos on columns types and numer of null values  
    tab_info = pd.DataFrame(df_initial.dtypes).T.rename(index={0: 'column  
type'})  
    tab_info =  
tab_info.append(pd.DataFrame(df_initial.isnull().sum()).T.rename(index={0:  
'null values (nb)'}))  
    tab_info = tab_info.append(pd.DataFrame(df_initial.isnull().sum() /  
df_initial.shape[0] * 100).T.  
                                rename(index={0: 'null values (%)'}))  
    return tab_info
```

In [14]:

DfInfo(df)

Out[14]:

Closing Value

| | |
|-------------------------|---------|
| column type | float64 |
| null values (nb) | 0 |
| null values (%) | 0.0 |

In [15]:

df.index

Out[15]:

```
DatetimeIndex(['2000-01-04', '2000-01-05', '2000-01-06', '2000-01-07',  
               '2000-01-10', '2000-01-11', '2000-01-12', '2000-01-13',  
               '2000-01-14', '2000-01-18',  
               ...  
               '2018-06-26', '2018-06-27', '2018-06-28', '2018-06-29',  
               '2018-07-02', '2018-07-03', '2018-07-04', '2018-07-05',  
               '2018-07-06', '2018-07-09'],  
              dtype='datetime64[ns]', name='Date', length=4673, freq=None)
```

In [16]:

```
y = df['Closing Value'].resample('MS').mean()
```

In [17]:

```
y.plot(figsize=(15, 6))
plt.show()
```

In [18]:

```
rcParams['figure.figsize'] = 18, 8
decomposition = sm.tsa.seasonal_decompose(y, model='additive')
fig = decomposition.plot()
plt.show()
```

In [19]:

```
sc = MinMaxScaler(feature_range = (0, 1))
df = sc.fit_transform(df)
```

TRAINING AND TESTING

In [20]:

```
train_size = int(len(df) * 0.70)
test_size = len(df) - train_size
train, test = df[0:train_size, :], df[train_size:len(df), :]
```

In [21]:

```
def create_data_set(_data_set, _look_back=1):
    data_x, data_y = [], []
    for i in range(len(_data_set) - _look_back - 1):
        a = _data_set[i:(i + _look_back), 0]
        data_x.append(a)
        data_y.append(_data_set[i + _look_back, 0])
    return np.array(data_x), np.array(data_y)
```

In [23]:

```
look_back = 90
X_train, Y_train, X_test, Y_test = [], [], [], []
X_train, Y_train = create_data_set(train, look_back)
X_train = np.reshape(X_train, (X_train.shape[0], X_train.shape[1], 1))
X_test, Y_test = create_data_set(test, look_back)
X_test = np.reshape(X_test, (X_test.shape[0], X_test.shape[1], 1))
```

LSTM LAYER

In [24]:

```
regressor = Sequential()

regressor.add(LSTM(units = 60, return_sequences = True, input_shape =
(X_train.shape[1], 1)))
regressor.add(Dropout(0.1))

regressor.add(LSTM(units = 60, return_sequences = True))
regressor.add(Dropout(0.1))

regressor.add(LSTM(units = 60))
regressor.add(Dropout(0.1))

regressor.add(Dense(units = 1))
```

```

regressor.compile(optimizer = 'adam', loss = 'mean_squared_error')
reduce_lr = ReduceLROnPlateau(monitor='val_loss',patience=5)
history =regressor.fit(X_train, Y_train, epochs = 20, batch_size =
15,validation_data=(X_test, Y_test), callbacks=[reduce_lr],shuffle=False)

Epoch 1/20
212/212 [=====] - 22s 85ms/step - loss: 0.0045 - val
_loss: 0.0321 - lr: 0.0010
Epoch 2/20
212/212 [=====] - 17s 79ms/step - loss: 0.0122 - val
_loss: 0.0445 - lr: 0.0010
Epoch 3/20
212/212 [=====] - 17s 81ms/step - loss: 0.0110 - val
_loss: 0.0513 - lr: 0.0010
Epoch 4/20
212/212 [=====] - 17s 78ms/step - loss: 0.0164 - val
_loss: 0.0483 - lr: 0.0010
Epoch 5/20
212/212 [=====] - 17s 79ms/step - loss: 0.0187 - val
_loss: 0.0549 - lr: 0.0010
Epoch 6/20
212/212 [=====] - 17s 79ms/step - loss: 0.0185 - val
_loss: 0.0478 - lr: 0.0010
Epoch 7/20
212/212 [=====] - 17s 80ms/step - loss: 0.0186 - val
_loss: 0.0031 - lr: 1.0000e-04
Epoch 8/20
212/212 [=====] - 17s 80ms/step - loss: 0.0032 - val
_loss: 0.0024 - lr: 1.0000e-04
Epoch 9/20
212/212 [=====] - 17s 79ms/step - loss: 0.0024 - val
_loss: 0.0019 - lr: 1.0000e-04
Epoch 10/20
212/212 [=====] - 17s 79ms/step - loss: 0.0021 - val
_loss: 0.0017 - lr: 1.0000e-04
Epoch 11/20
212/212 [=====] - 17s 82ms/step - loss: 0.0016 - val
_loss: 0.0017 - lr: 1.0000e-04
Epoch 12/20
212/212 [=====] - 17s 79ms/step - loss: 0.0015 - val
_loss: 0.0016 - lr: 1.0000e-04
Epoch 13/20
212/212 [=====] - 17s 80ms/step - loss: 0.0013 - val
_loss: 0.0014 - lr: 1.0000e-04
Epoch 14/20
212/212 [=====] - 17s 79ms/step - loss: 0.0012 - val
_loss: 0.0013 - lr: 1.0000e-04
Epoch 15/20
212/212 [=====] - 17s 78ms/step - loss: 0.0011 - val
_loss: 0.0013 - lr: 1.0000e-04
Epoch 16/20

```

```

212/212 [=====] - 17s 78ms/step - loss: 9.6755e-04 -
val_loss: 0.0013 - lr: 1.0000e-04
Epoch 17/20
212/212 [=====] - 17s 80ms/step - loss: 0.0010 - val
_loss: 0.0013 - lr: 1.0000e-04
Epoch 18/20
212/212 [=====] - 17s 79ms/step - loss: 9.4863e-04 -
val_loss: 0.0013 - lr: 1.0000e-04
Epoch 19/20
212/212 [=====] - 17s 81ms/step - loss: 9.1643e-04 -
val_loss: 0.0014 - lr: 1.0000e-04
Epoch 20/20
212/212 [=====] - 18s 84ms/step - loss: 9.3516e-04 -
val_loss: 0.0013 - lr: 1.0000e-04

```

MODEL TRAINING

In [25]:

```

train_predict = regressor.predict(X_train)
test_predict = regressor.predict(X_test)

100/100 [=====] - 4s 27ms/step
41/41 [=====] - 1s 28ms/step

```

In [26]:

```

train_predict = sc.inverse_transform(train_predict)
Y_train = sc.inverse_transform([Y_train])
test_predict = sc.inverse_transform(test_predict)
Y_test = sc.inverse_transform([Y_test])

```

PREDICTION

In [27]:

```

print('Train Mean Absolute Error:', mean_absolute_error(Y_train[0],
train_predict[:,0]))
print('Train Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_train[0],
train_predict[:,0])))
print('Test Mean Absolute Error:', mean_absolute_error(Y_test[0],
test_predict[:,0]))
print('Test Root Mean Squared Error:', np.sqrt(mean_squared_error(Y_test[0],
test_predict[:,0])))
plt.figure(figsize=(8,4))
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Test Loss')
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epochs')
plt.legend(loc='upper right')
plt.show();

```

```

Train Mean Absolute Error: 2.7544575163643317
Train Root Mean Squared Error: 3.5074279586248873
Test Mean Absolute Error: 2.373878536093426
Test Root Mean Squared Error: 5.286033010892939

```

In [28]:

```

aa=[x for x in range(180)]

```

```
plt.figure(figsize=(8,4))
plt.plot(aa, Y_test[0][:180], marker='.', label="actual")
plt.plot(aa, test_predict[:,0][:180], 'r', label="prediction")
plt.tight_layout()
sns.despine(top=True)
plt.subplots_adjust(left=0.07)
plt.ylabel('Price', size=15)
plt.xlabel('Time step', size=15)
plt.legend(fontsize=15)
plt.show();
```